

MODELS

LINEAR REGRESSION:

```
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error


# Load the dataset using pandas
data = pd.read_csv('your_dataset.csv') # Replace 'your_dataset.csv' with the path to your dataset


# Assuming your dataset has columns 'X1', 'X2', ..., 'Xn' for features and 'y' for target variable
X = data.drop('y', axis=1) # Input features (excluding the target variable)
y = data['y']             # Target variable


# Split the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)


# Standardize the features using StandardScaler (although not necessary for Linear Regression)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)


# Create a Linear Regression model
model = LinearRegression()


# Fit the model to the training data
model.fit(X_train_scaled, y_train)
```

```
# Predict the target variable for the test set
```

```
y_pred = model.predict(X_test_scaled)
```

```
# Calculate the mean squared error
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
print("Mean Squared Error:", mse)
```

LOGISTIC REGRESSION:

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.metrics import accuracy_score
```

```
# Load the dataset using pandas
```

```
data = pd.read_csv('your_dataset.csv') # Replace 'your_dataset.csv' with the path to your dataset
```

```
# Assuming your dataset has columns 'X1', 'X2', ..., 'Xn' for features and 'y' for target variable
```

```
X = data.drop('y', axis=1) # Input features (excluding the target variable)
```

```
y = data['y'] # Target variable
```

```
# Split the dataset into training and testing sets (80% train, 20% test)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Standardize the features using StandardScaler
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```

# Create a Logistic Regression model
model = LogisticRegression()

# Fit the model to the training data
model.fit(X_train_scaled, y_train)

# Predict the labels for the test set
y_pred = model.predict(X_test_scaled)

# Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

```

NAÏVE BAYES:

```

import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score

# Load the dataset using pandas
data = pd.read_csv('your_dataset.csv') # Replace 'your_dataset.csv' with the path to your dataset

# Assuming your dataset has columns 'X1', 'X2', ..., 'Xn' for features and 'y' for target variable
X = data.drop('y', axis=1) # Input features (excluding the target variable)
y = data['y']              # Target variable

# Split the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

```
# Standardize the features using StandardScaler
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
# Create a Gaussian Naive Bayes model
```

```
model = GaussianNB()
```

```
# Fit the model to the training data
```

```
model.fit(X_train_scaled, y_train)
```

```
# Predict the labels for the test set
```

```
y_pred = model.predict(X_test_scaled)
```

```
# Calculate the accuracy of the model
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print("Accuracy:", accuracy)
```

KNN:

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.metrics import accuracy_score
```

```
# Load the dataset using pandas
```

```
data = pd.read_csv('your_dataset.csv') # Replace 'your_dataset.csv' with the path to your dataset
```

```
# Assuming your dataset has columns 'X1', 'X2', ..., 'Xn' for features and 'y' for target variable
```

```
X = data.drop('y', axis=1) # Input features (excluding the target variable)
```

```

y = data['y']          # Target variable

# Split the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize the features using StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Create a K-Nearest Neighbors model with k=5
model = KNeighborsClassifier(n_neighbors=5)

# Fit the model to the training data
model.fit(X_train_scaled, y_train)

# Predict the labels for the test set
y_pred = model.predict(X_test_scaled)

# Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

```

HIERARCHIAL CLUSTERING:

```

import pandas as pd

from sklearn.preprocessing import StandardScaler
from sklearn.cluster import AgglomerativeClustering

# Load the dataset using pandas
data = pd.read_csv('your_dataset.csv') # Replace 'your_dataset.csv' with the path to your dataset

```

```
# Assuming your dataset has columns 'X1', 'X2', ..., 'Xn' for features
```

```
X = data[['X1', 'X2', ..., 'Xn']] # Input features
```

```
# Standardize the features using StandardScaler
```

```
scaler = StandardScaler()
```

```
X_scaled = scaler.fit_transform(X)
```

```
# Create a Hierarchical clustering model
```

```
model = AgglomerativeClustering(n_clusters=3) # Example: 3 clusters
```

```
# Fit the model to the data
```

```
model.fit(X_scaled)
```

```
# Get cluster labels
```

```
labels = model.labels_
```

BOOSTED TREE:

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.ensemble import GradientBoostingClassifier
```

```
from sklearn.metrics import accuracy_score
```

```
# Load the dataset using pandas
```

```
data = pd.read_csv('your_dataset.csv') # Replace 'your_dataset.csv' with the path to your dataset
```

```
# Assuming your dataset has columns 'X1', 'X2', ..., 'Xn' for features and 'y' for target variable
```

```
X = data[['X1', 'X2', ..., 'Xn']] # Input features
```

```
y = data['y'] # Target variable
```

```
# Split the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize the features using StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Create a Gradient Boosting Classifier model
model = GradientBoostingClassifier()

# Fit the model to the training data
model.fit(X_train_scaled, y_train)

# Predict the labels for the test set
y_pred = model.predict(X_test_scaled)

# Calculate the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

XGBOOSTER TREE:

```
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score

# Load the dataset using pandas
```

```
data = pd.read_csv('your_dataset.csv') # Replace 'your_dataset.csv' with the path to your dataset
```

```
# Assuming your dataset has columns 'X1', 'X2', ..., 'Xn' for features and 'y' for target variable
```

```
X = data[['X1', 'X2', ..., 'Xn']] # Input features
```

```
y = data['y'] # Target variable
```

```
# Split the dataset into training and testing sets (80% train, 20% test)
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Standardize the features using StandardScaler
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
# Create an XGBoost Classifier model
```

```
model = XGBClassifier()
```

```
# Fit the model to the training data
```

```
model.fit(X_train_scaled, y_train)
```

```
# Predict the labels for the test set
```

```
y_pred = model.predict(X_test_scaled)
```

```
# Calculate the accuracy of the model
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print("Accuracy:", accuracy)
```

BAG OF WORDS DOCUMENT CLASSIFIER[EMAIL SPAM]:

```
import pandas as pd
```

```
from sklearn.feature_extraction.text import CountVectorizer
```



```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.naive_bayes import MultinomialNB
from sklearn.metrics import accuracy_score, confusion_matrix

# Load the dataset
data = pd.read_csv('spam_dataset.csv') # Replace 'spam_dataset.csv' with your dataset file

# Assuming your dataset has columns 'text' for email content and 'label' for spam/ham
X_text = data['text'] # Email content
y = data['label'] # Spam/ham label

# Convert text data into numerical features using bag-of-words approach
vectorizer = CountVectorizer()
X_bow = vectorizer.fit_transform(X_text)

# Split the dataset into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X_bow, y, test_size=0.2, random_state=42)

# Convert to dense format for StandardScaler
X_train_dense = X_train.toarray()
X_test_dense = X_test.toarray()

# Standardize the features using StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_dense)
X_test_scaled = scaler.transform(X_test_dense)

# Train a Naive Bayes classifier
```

```
classifier = MultinomialNB()
classifier.fit(X_train_scaled, y_train)

# Predict on the test set
y_pred = classifier.predict(X_test_scaled)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

# Confusion matrix
conf_matrix = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(conf_matrix)
```

ENSEMBLE MODELS[RANDOM FOREST]:

```
import pandas as pd

from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score

# Load dataset
df = pd.read_csv('your_dataset.csv')

# Assuming the target variable is in a column named 'target'
X = df.drop(columns=['target']) # Features
y = df['target'] # Target variable

# Split dataset into training set and test set
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
# Standardize features
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
# Create Random Forest classifier object
```

```
clf = RandomForestClassifier()
```

```
# Train Random Forest Classifier
```

```
clf = clf.fit(X_train_scaled, y_train)
```

```
# Predict the response for test dataset
```

```
y_pred = clf.predict(X_test_scaled)
```

```
# Model Accuracy
```

```
print("Accuracy:", accuracy_score(y_test, y_pred))
```

PCA:

```
# Importing necessary libraries
```

```
import pandas as pd
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.decomposition import PCA
```

```
import matplotlib.pyplot as plt
```

```
# Load your dataset
```

```
df = pd.read_csv('your_dataset.csv')
```

```
# Assuming your dataset has features and possibly a target variable
```

```

# Separate features from the target variable (if applicable)

X = df.drop(columns=['target_column_name']) # Adjust 'target_column_name' with your target variable
name

y = df['target_column_name'] # Adjust 'target_column_name' with your target variable name if
applicable


# Standardize the features

scaler = StandardScaler()

X_scaled = scaler.fit_transform(X)


# Perform PCA

pca = PCA(n_components=2)

X_pca = pca.fit_transform(X_scaled)


# Visualize the PCA results

plt.figure(figsize=(8, 6))

plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y, cmap='viridis', edgecolor='k') # Assuming y is available, else
remove c=y

plt.xlabel('Principal Component 1')

plt.ylabel('Principal Component 2')

plt.title('PCA of Your Dataset with StandardScaler')

plt.colorbar(label='Target') # Adjust if you have a target variable

plt.show()

```

SVC:

```

# Importing necessary libraries

import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.preprocessing import StandardScaler

from sklearn.svm import SVC

from sklearn.metrics import accuracy_score

```

```
# Load your dataset
```

```
df = pd.read_csv('your_dataset.csv')
```

```
# Assuming your dataset has features and a target variable
```

```
# Separate features from the target variable
```

```
X = df.drop(columns=['target_column_name']) # Adjust 'target_column_name' with your target variable name
```

```
y = df['target_column_name'] # Adjust 'target_column_name' with your target variable name
```

```
# Split the dataset into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Standardize the features
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
# Initialize the SVC classifier
```

```
svc = SVC()
```

```
# Train the SVC classifier
```

```
svc.fit(X_train_scaled, y_train)
```

```
# Predict on the test set
```

```
y_pred = svc.predict(X_test_scaled)
```

```
# Calculate accuracy
```

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print(f"Accuracy: {accuracy}")
```

SVR:

```
# Importing necessary libraries
```

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import StandardScaler
```

```
from sklearn.svm import SVR
```

```
from sklearn.metrics import mean_squared_error
```

```
# Load your dataset
```

```
df = pd.read_csv('your_dataset.csv')
```

```
# Assuming your dataset has features and a target variable
```

```
# Separate features from the target variable
```

```
X = df.drop(columns=['target_column_name']) # Adjust 'target_column_name' with your target variable name
```

```
y = df['target_column_name'] # Adjust 'target_column_name' with your target variable name
```

```
# Split the dataset into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Standardize the features
```

```
scaler = StandardScaler()
```

```
X_train_scaled = scaler.fit_transform(X_train)
```

```
X_test_scaled = scaler.transform(X_test)
```

```
# Initialize the SVR model
```

```
svr = SVR()
```

```
# Train the SVR model
```

```
svr.fit(X_train_scaled, y_train)
```

```
# Predict on the test set
```

```
y_pred = svr.predict(X_test_scaled)
```

```
# Calculate Mean Squared Error
```

```
mse = mean_squared_error(y_test, y_pred)
```

```
print(f"Mean Squared Error: {mse}")
```